

Insurgents Speak

Data Dictionary

Version 0.2 · Draft

Cem Emrence

`emrence@umich.edu`

This document describes a corpus in development.
Field definitions and codes may change before the first release.

Overview

Insurgents Speak is a multilingual corpus of texts produced by insurgent organizations, movements, and other political outsiders, paired with conflict data so that what these groups said can be studied alongside what they did.


The corpus is distributed in three tiers. **Metadata** is public for every document located, whether or not its text can be redistributed. **Derived features**—token counts, part-of-speech distributions, dependency relations, and summary measures such as verb density—are public for every document processed, since these are not reproductions of the underlying text. **Full text** is distributed where rights permit. The design follows the non-consumptive research model established by the HathiTrust Research Center, which allows scholars to compute on in-copyright material without receiving it.

Groups that appear in the Uppsala Conflict Data Program carry a `ucdp_id`, which allows documents to be merged with conflict event and actor data. Groups outside UCDP have a `group_id` but no `ucdp_id`: the corpus is not limited to UCDP's universe, and is designed to grow toward a definition of its population built on an actor's position relative to institutional politics rather than on the intensity of its violence.

Structure

The corpus is distributed as five linked tables.

Table	Contents	Key
group	One row per organization	<code>group_id</code>
relations	One row per relation between organizations	<code>relation_id</code>
document	One row per document	<code>doc_id</code>
features	Derived measures per processed document	<code>doc_id</code>
contributor	One row per contributor	<code>contributor_id</code>



The relations table references the group table twice, through `group_id_from` and `group_id_to`.

Group table

One row per organization.

Field	Description
group_id	Identifier for the organization, assigned by the project. Every group has one, including groups that appear in no existing conflict dataset.
group_short	The organization's own acronym, reproduced as the group uses it. Acronyms might not be unique across countries and periods, so use group_id to identify groups.
group_name	Full name of the organization, in the language the group uses.
group_name_en	Name used for the organization in English-language scholarship. Blank only where no English form is established.
ucdp_id	Actor identifier from the Uppsala Conflict Data Program, where the group appears there. Blank otherwise.
ucdp_note	Notes on the mapping to UCDP where it is imperfect: coverage that begins later than the organization's activity, name changes across UCDP's coding, or cases where UCDP's actor definition does not align with the organization as it appears in the corpus.
country_conflict	ISO 3166 alpha-3 code for the country of the conflict the organization is party to.

Relations table

One row per relation between organizations. A group may appear in several rows.

Field	Description
relation_id	Identifier for the relation.
group_id_from	The earlier or originating organization.
group_id_to	The later or resulting organization.
relation_type	Type of relation, coded numerically; see below.
year	Year the relation took effect, where known.
contested	Whether the relation is disputed by the parties or in the scholarship. Splits in particular are often claimed differently by each faction.
note	Brief description, where the relation needs explanation.

relation_type codes

Code	Value	Description
1	Split	group_id_to broke away from group_id_from
2	Merger	group_id_from combined into group_id_to; one row per constituent group
3	Rename	Same organization, new name
4	Successor	group_id_to claims continuity with a dissolved group_id_from
5	Front	group_id_to operates as an affiliated or subordinate body

Renames apply where the organization's continuity is uncontested; splits where a new organization emerges alongside or in place of the original.

Document table

One row per document.

Field	Description
doc_id	Unique identifier for the document, assigned in sequence as documents enter the corpus. The number reflects processing order and encodes nothing about the group, date, or type, which keeps identifiers stable when documents are later recoded, redated, or added.
group_id	Links to the group table.
year	Year of publication as given by the document.
month	Month of publication, 1–12. Left blank when the document gives only a year.
day	Day of publication, where the document gives an exact date. Populated mainly for communiqués and other dated standalone documents.
language	ISO 639-3 code for the language of the text itself.
doc_type	Type of document, coded numerically; see below.
country_publication	ISO 3166 alpha-3 code for where the document was published. Differs from <code>country_conflict</code> for organizations publishing in exile, which is common in this population.
publication_title	Title of the periodical or series. Blank for standalone documents.
volume_no	Volume number as printed, where the publication uses one.
issue_no	Issue number as printed.
pages	Number of pages in the document.
text_source	Whether the text is born-digital, OCR'd from images, or transcribed by hand.
ocr_engine	Software and version used, where the text was OCR'd.
source_archive	Repository or collection holding the document.
source_url	Link to the source, where one exists.
access_date	Date the document was retrieved.
rights_status	Basis on which the document's text is distributed; see below. Metadata and derived features are available for all documents regardless of this value.
text_file	Name of the file containing the full text, keyed to <code>doc_id</code> . Blank where <code>rights_status</code> is restricted.
contributor_id	Links to the contributor table.

doc_type codes

Code	Value	Description
1	Periodical	Numbered issue of a recurring publication
2	Communiqué	Short official statement announcing or explaining a specific event, decision, or position
3	Manifesto	Programmatic or founding document
4	Leaflet	Single-sheet document distributed directly to the community
5	Transcript	Broadcast or speech, transcribed
6	Web	Born-digital posting
7	Book	Published under organizational or leadership authorship
8	Interview	Interview with movement figures, published or transcribed
9	Pamphlet	Short publication produced by or alongside a periodical
10	Proceedings	Record of a congress, conference, or other organizational meeting
11	Legal	Trial transcripts, defense statements, and other court material, where produced or published by the organization

Where a document could fall under more than one type, the more specific applies: a defense statement delivered in court is coded *Legal* rather than *Transcript*.

Document length varies widely across types: a book and a leaflet each count as one document. Analyses computed per document should account for this, using `pages` or `n_tokens`.

rights_status values

Value	Description
public-domain	Copyright expired or never attached
licensed	Distributed under an open license
permission	Rightsholder granted use
restricted	Text not redistributable; metadata and features still available

Features table

Derived measures, available for every processed document including those whose text cannot be redistributed. Joined to the document table on `doc_id`.

Field	Description
<code>doc_id</code>	Links to the document table.
<code>n_tokens</code>	Total tokens in the document.
<code>n_sentences</code>	Total sentences.
<code>verb_pct</code>	Verbs as a percentage of all tagged tokens.
<code>noun_pct</code> , <code>adj_pct</code> , <code>adv_pct</code> , <code>pron_pct</code>	Corresponding shares for other major parts of speech.
<code>ttr</code>	Type-token ratio, a measure of lexical variety.
<code>mean_sentence_length</code>	Mean tokens per sentence.

Field	Description
pos_full	Counts for all seventeen Universal Dependencies part-of-speech classes.
dep_full	Counts for dependency relations.
tagger	Software and model used, for example <code>udpipe 0.8 / tr-imst</code> . Feature comparability across languages depends on which model produced them.

Features are summary statistics rather than reproductions; the underlying text cannot be reconstructed from them.

Contributor table

One row per contributor.

Field	Description
contributor_id	Identifier assigned by the project.
contributor_name	Name, recorded only where the contributor consents to attribution.
contributor_orcid	ORCID identifier, where the contributor provides one.
affiliation	Institutional affiliation, where given.
n_documents	Number of validated documents submitted.

Coding conventions

Names. Organization names carry political claims. `group_short` and `group_name` reproduce the form the organization uses for itself, and are not standardized, translated, or abbreviated further.

Language. `language` records the language of the text, not the language implied by a publication title or by the group's claimed constituency.

Dates. A blank `month` or `day` records the precision the document itself provides.

Excerpts. Where an excerpt is shown, it runs up to 250 characters from the middle of the document, beginning and ending at sentence boundaries. Excerpts are lightly corrected for OCR error so that they read as the source does; the distributed text is not corrected.

Citation

Emrence, Cem. *Insurgents Speak: A Multilingual Corpus of Insurgent Texts*. Version 0.2.