

Insurgents Speak

Data Dictionary

Version 0.3 · Draft

Cem Emrence

`emrence@umich.edu`

This document describes a corpus in development.
Field definitions and codes may change before the first release.

Overview

Insurgents Speak is a multilingual corpus of texts produced by insurgent organizations, paired with conflict data so that what these groups said can be studied alongside what they did.

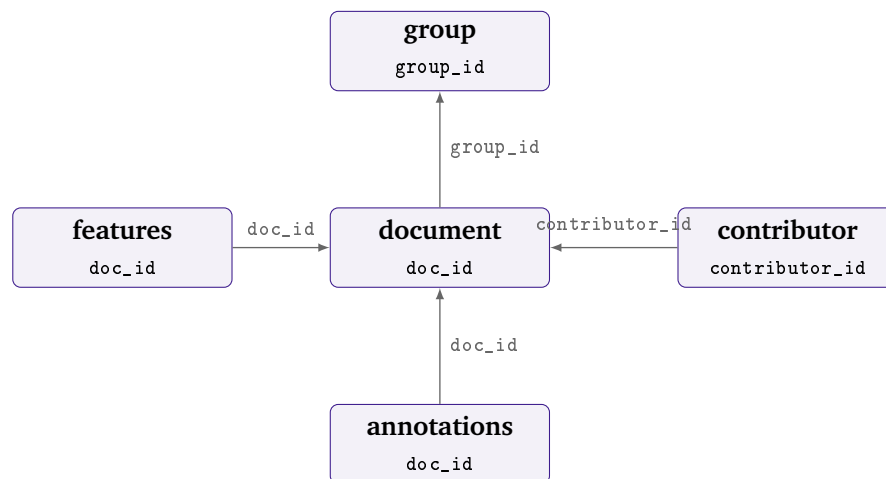
The corpus is distributed in three tiers. **Metadata** is public for every document located, whether or not its text can be redistributed. **Derived features**—token counts, part-of-speech distributions, dependency relations, and summary measures such as verb density—are public for every document processed, since these are not reproductions of the underlying text. **Full text** is distributed where rights permit. This follows the non-consumptive model established by the HathiTrust Research Center, which lets scholars compute on in-copyright material without receiving it.

Groups that appear in the Uppsala Conflict Data Program carry a `ucdp_id`, which allows documents to be merged with conflict event and actor data. Groups outside UCDP have a `group_id` but no `ucdp_id`. Hence, the corpus will not be limited to UCDP's universe.

Structure

The corpus is distributed as five linked tables.

Table	Contents	Key
group	One row per organization	group_id
document	One row per document	doc_id
features	Document-level measures	doc_id
annotations	One row per document per tag	doc_id
contributor	One row per contributor	contributor_id



At the beta release, the group, document, features, and annotations tables are populated. The contributor table fills as moderated crowdsourcing begins.

Group table

One row per organization.

Field	Description
group_id	Identifier for the organization, assigned by the project. Every group has one, including groups that appear in no existing conflict dataset.
group_short	The organization's own acronym, reproduced as the group uses it.
group_name	Full name of the organization, in the language the group uses.
group_name_en	Name used for the organization in English-language scholarship.
ucdp_id	Actor identifier from the Uppsala Conflict Data Program, where the group appears there. Blank otherwise.
ucdp_note	Notes on when mapping to UCDP is imperfect (coverage, name changes, definitional mismatch).
country_conflict	ISO 3166 alpha-3 code for the country of the conflict the organization is party to.

Document table

One row per document.

Field	Description
doc_id	Unique identifier for the document, assigned in sequence as documents enter the corpus.
group_id	Links to the group table.
year	Year of publication as given by the document.
month	Month of publication, 1–12. Left blank when the document gives only a year.
day	Day of publication, where the document gives an exact date. Populated mainly for communiqués and other dated standalone documents.
language	ISO 639-3 code for the language of text.
doc_type	Type of document, coded numerically; see below.
country_publication	ISO 3166 alpha-3 code for where the document was published. Differs from country_conflict for organizations publishing abroad.
publication_title	Title of the periodical or series. Blank for standalone documents.
volume_no	Volume number as printed, where the publication uses one.
issue_no	Issue number as printed.
pages	Number of pages in the document.
text_source	How the text was produced: born-digital, OCR, or transcription.
ocr_engine	Software and version used, where the text was OCR'd.
source_archive	Repository or collection holding the document.
source_url	Link to the source, where one exists.
access_date	Date the document was retrieved.

Field	Description
rights_status	Basis on which the document's text is distributed; see below. Metadata and derived features are available for all documents regardless of this value.
text_file	Name of the file containing the full text, doc_id zero-padded to six digits. Blank where rights_status is restricted.
contributor_id	Links to the contributor table.

doc_type codes

Code	Value	Description
1	Periodical	Numbered issue of a recurring publication
2	Communiqué	Short official statement announcing or explaining a specific event, decision, or position
3	Manifesto	Programmatic or founding document
4	Leaflet	Single-sheet document distributed directly to the community
5	Transcript	Broadcast or speech, transcribed
6	Web	Born-digital posting
7	Book	Published under organizational or leadership authorship
8	Interview	Interview with movement figures, published or transcribed
9	Pamphlet	Short publication produced by or alongside a periodical
10	Proceedings	Record of a congress, conference, or other organizational meeting
11	Legal	Trial transcripts, defense statements, and other court material, where produced or published by the organization

Where a document could fall under more than one type, the more specific applies: a defense statement delivered in court is coded *Legal* rather than *Transcript*.

rights_status values

Value	Description
public-domain	Copyright expired or never attached
licensed	Distributed under an open license
permission	Rightsholder granted use
restricted	Text not redistributable; metadata and features still available

Features table

Document-level measures, available for every processed document including those whose text cannot be redistributed. Joined to the document table on doc_id.

Field	Description
doc_id	Links to the document table.
n_tokens	Total tokens in the document.

Field	Description
n_sentences	Total sentences.
mean_sentence_length	Mean tokens per sentence.
mattr	Moving-average type-token ratio over a 500-token window, a measure of lexical variety that is comparable across documents of different lengths.
tagger	Software and version used, for example <code>udpipe 0.8</code> .
treebank	Universal Dependencies treebank the model was trained on.

Features are summary statistics rather than reproductions; the underlying text cannot be reconstructed from them.

Annotations table

One row per document per tag.

Field	Description
doc_id	Links to the document table.
layer	Annotation layer: <code>pos</code> for part-of-speech tags, <code>dep</code> for dependency relations.
tag	Universal Dependencies tag or relation.
n	Count in that document.

Divide `n` by `n_tokens` from the features table for shares.

Contributor table

One row per contributor.

Field	Description
contributor_id	Identifier assigned by the project.
contributor_name	Name, recorded only where the contributor consents to attribution.
contributor_orcid	ORCID identifier, where the contributor provides one.
affiliation	Institutional affiliation, where given.
n_documents	Number of validated documents submitted.

Citation

Emrence, Cem. *Insurgents Speak: A Multilingual Corpus of Insurgent Texts*. Version 0.3.